

DDPO training at Re=2000 — reward, raw component distances, and the GT probe (curves smoothed over 9 iterations)

